

# Neural Networks & Deep Learning – ICP 8

Name: Vyshnavi Bussa

Student ID: 700757313

GitHub Link: <https://github.com/Vyshnavi1157/NN-ICP8.git>

Video Link: <https://drive.google.com/file/d/1VZ-0oI43-Co-XsxZzGmezO4zoHhuIowC/view?usp=sharing>

```
In [8]: import keras
        from keras.models import Sequential
        from keras.preprocessing import image
        from keras.layers import Activation,Dense,Dropout,Conv2D,Flatten,MaxPooling2D,BatchNormalization
        from keras.datasets import cifar10
        from keras import optimizers
        from matplotlib import pyplot as plt
```


```
In [10]: #generate cifar10 data
        (x_train,y_train),(x_test,y_test) = cifar10.load_data()
```

```
In [9]: #config parameters
        num_classes = 10
        input_shape = x_train.shape[1:4]
        optimizer = optimizers.Adam(lr=0.001)
```


```
In [11]: #convert label to one-hot
        one_hot_y_train = keras.utils.to_categorical(y_train,num_classes=num_classes)
        one_hot_y_test = keras.utils.to_categorical(y_test,num_classes=num_classes)
```

```
In [12]: # check data
        plt.imshow(x_train[1])
        print(x_train[1].shape)
```

(32, 32, 3)



(32, 32, 3)



In [13]:

```
# build model(similar to VGG16, only change the input and output shape)
model = Sequential()
model.add(Conv2D(64,(3,3),activation='relu',input_shape=input_shape,padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(64,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))
```

```

model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(4096,activation='relu'))
model.add(Dense(2048,activation='relu'))
model.add(Dense(1024,activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes))
model.add(Activation('softmax'))

```

```
In [14]: model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
```

```
In [14]: model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
```

```
In [15]: model.summary()
```

```

Model: "sequential_1"

```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 64)	1792
batch_normalization (Batch Normalization)	(None, 32, 32, 64)	256
conv2d_2 (Conv2D)	(None, 32, 32, 64)	36928
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 64)	256
max_pooling2d (MaxPooling2D)	(None, 16, 16, 64)	0
dropout (Dropout)	(None, 16, 16, 64)	0

```
In [ ]: history = model.fit(x=x_train, y=one_hot_y_train, batch_size=128, epochs=30, validation_split=0.1)
```

```

Epoch 1/30
352/352 [=====] - 29s 81ms/step - loss: 1.7083 - accuracy: 0.3346 - val_loss: 2.6400 - val_accuracy: 0.3686
Epoch 2/30
352/352 [=====] - 27s 76ms/step - loss: 1.3159 - accuracy: 0.5254 - val_loss: 1.6935 - val_accuracy: 0.5590
Epoch 3/30
352/352 [=====] - 27s 77ms/step - loss: 1.0110 - accuracy: 0.6538 - val_loss: 1.0628 - val_accuracy: 0.6556
Epoch 4/30
352/352 [=====] - 27s 78ms/step - loss: 0.8493 - accuracy: 0.7144 - val_loss: 0.9995 - val_accuracy: 0.6856
Epoch 5/30
352/352 [=====] - 28s 79ms/step - loss: 0.7343 - accuracy: 0.7566 - val_loss: 0.8316 - val_accuracy: 0.7326
Epoch 6/30
352/352 [=====] - 28s 79ms/step - loss: 0.6515 - accuracy: 0.7907 - val_loss: 0.8145 - val_accuracy: 0.7326

```

```

0.8430
Epoch 22/30
352/352 [=====] - 28s 78ms/step - loss: 0.1242 - accuracy: 0.9620 - val_loss: 0.6557 - val_accuracy:
0.8544
Epoch 23/30
352/352 [=====] - 28s 80ms/step - loss: 0.1144 - accuracy: 0.9639 - val_loss: 0.7322 - val_accuracy:
0.8400
Epoch 24/30
352/352 [=====] - 28s 80ms/step - loss: 0.1086 - accuracy: 0.9653 - val_loss: 0.6893 - val_accuracy:
0.8548
Epoch 25/30
352/352 [=====] - 28s 78ms/step - loss: 0.1313 - accuracy: 0.9619 - val_loss: 0.6570 - val_accuracy:
0.8412
Epoch 26/30
352/352 [=====] - 28s 80ms/step - loss: 0.1069 - accuracy: 0.9673 - val_loss: 0.6935 - val_accuracy:
0.8454
Epoch 27/30
352/352 [=====] - 28s 78ms/step - loss: 0.0962 - accuracy: 0.9703 - val_loss: 0.7319 - val_accuracy:
0.8376
Epoch 28/30
352/352 [=====] - 28s 78ms/step - loss: 0.0871 - accuracy: 0.9728 - val_loss: 0.7403 - val_accuracy:
0.8480
Epoch 29/30
352/352 [=====] - 28s 79ms/step - loss: 0.0811 - accuracy: 0.9749 - val_loss: 0.6592 - val_accuracy:
0.8510
Epoch 30/30
352/352 [=====] - 28s 78ms/step - loss: 0.0753 - accuracy: 0.9770 - val_loss: 0.6631 - val_accuracy:
0.8642

```

```

In [ ]: # evaluate
print(model.metrics_names)
model.evaluate(x=x_test,y=one_hot_y_test,batch_size=512)

['loss', 'accuracy']
20/20 [=====] - 5s 132ms/step - loss: 0.6627 - accuracy: 0.8592

```

Out[21]: [0.6626988053321838, 0.8592000007629395]

```

In [ ]: model.save("keras-VGG16-cifar10.h5")
plt.imshow(x_test[1000])

result = model.predict(x_test[1000:1001]).tolist()
predict = 0


```

```
Out[21]: [0.6626988053321838, 0.8592000007629395]
```

```
In [ ]: model.save("keras-VGG16-cifar10.h5")
plt.imshow(x_test[1000])

result = model.predict(x_test[1000:1001]).tolist()
predict = 0
expect = y_test[1000][0]
for i,_ in enumerate(result[0]):
    if result[0][i] > result[0][predict]:
        predict = i
print("predict class:",predict)
print("expected class:",expect)


1/1 [=====] - 1s 740ms/step
predict class: 5
expected class: 5
```




0 5 10 15 20 25 30

```
In [16]: # save model
model.save("keras-VGG16-cifar10.h5")
```

```
In [ ]: #plot the training and validation accuracy
plt.plot(history.history['accuracy'], label='training accuracy')
plt.plot(history.history['val_accuracy'], label='validation accuracy')
plt.title('Training and validation accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



```
In [ ]: #plot the training and validation loss
plt.plot(history.history['loss'], label='training loss')
plt.plot(history.history['val_loss'], label='validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
In [ ]: import numpy as np
from sklearn.metrics import confusion_matrix

# calculate the confusion matrix
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
```

```

In [ ]: import numpy as np
        from sklearn.metrics import confusion_matrix


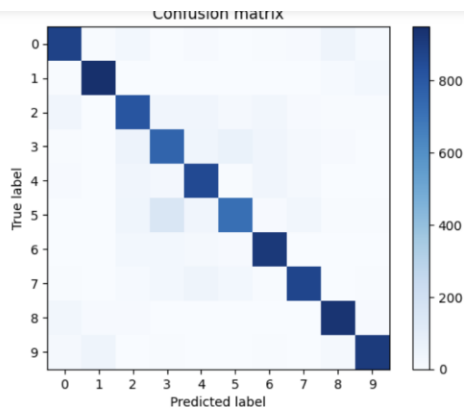
        # calculate the confusion matrix
        y_pred = model.predict(x_test)
        y_pred_classes = np.argmax(y_pred, axis=1)
        y_true = y_test.ravel()
        cm = confusion_matrix(y_true, y_pred_classes)

        # plot the confusion matrix
        plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
        plt.title('Confusion matrix')
        plt.colorbar()
        tick_marks = np.arange(num_classes)
        plt.xticks(tick_marks, range(num_classes))
        plt.yticks(tick_marks, range(num_classes))
        plt.xlabel('Predicted label')
        plt.ylabel('True label')
        plt.show()

        # plot a histogram of the predicted probabilities for a sample image
        plt.hist(y_pred[1000])
        plt.title('Predicted probabilities')
        plt.xlabel('Probability')
        plt.ylabel('Frequency')
        plt.show()

```

313/313 [=====] - 3s 8ms/step





#### Import libraries

```
In [11]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
from tensorflow.keras.datasets import mnist

from tensorflow.keras.optimizers import RMSprop
from keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, BatchNormalization

%matplotlib inline
```

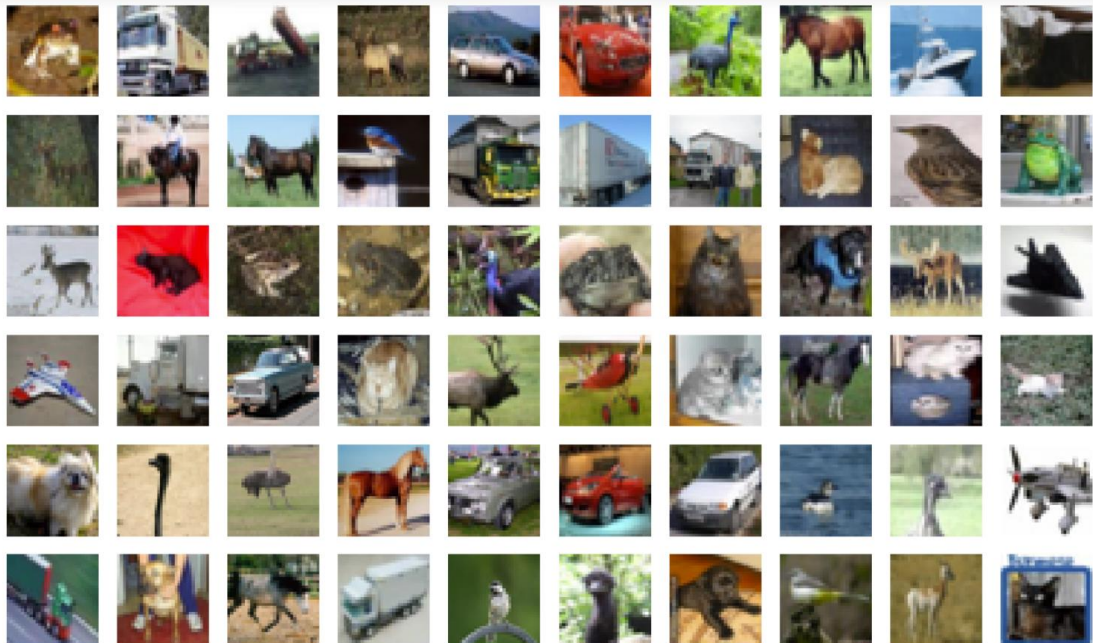
Extract data and train and test dataset


```
In [12]: #cifar100 = tf.keras.datasets.cifar100
(X_train,Y_train) , (X_test,Y_test) = cifar10.load_data()
```

```
In [13]: classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

Let's look into the dataset images

```
In [14]: plt.figure(figsize = (16,16))
for i in range(100):
    plt.subplot(10,10,i+1)
    plt.axis('off')
    plt.imshow(X_train[i], cmap = 'gray')
```





Training , Validating and Splitting trained and tested data

```
In [15]: from sklearn.model_selection import train_test_split
x_train, x_val, y_train, y_val = train_test_split(X_train, Y_train, test_size=0.2)
```

```
In [16]: from keras.utils.np_utils import to_categorical
y_train = to_categorical(y_train, num_classes = 10)
y_val = to_categorical(y_val, num_classes = 10)
```

```
In [17]: print(x_train.shape)
print(y_train.shape)
print(x_val.shape)
print(y_val.shape)
print(X_test.shape)
print(Y_test.shape)
```

```
(40000, 32, 32, 3)
(40000, 10)
(10000, 32, 32, 3)
(10000, 10)
(10000, 32, 32, 3)
(10000, 1)
```

```
In [18]: train_datagen = ImageDataGenerator(
    preprocessing_function = tf.keras.applications.vgg19.preprocess_input,
    rotation_range=10,
    zoom_range = 0.1,
    width_shift_range = 0.1,
    height_shift_range = 0.1,
    shear_range = 0.1,
    horizontal_flip = True
)
train_datagen.fit(x_train)

val_datagen = ImageDataGenerator(preprocessing_function = tf.keras.applications.vgg19.preprocess_input)
val_datagen.fit(x_val)
```

```
In [19]: from keras.callbacks import ReduceLROnPlateau
learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy',
    patience=3,
    verbose=1,
    factor=0.5,
    min_lr=0.00001)
```

We have used only 16 layers out of 19 layers in the CNN

```
In [20]: vgg_model = tf.keras.applications.VGG19(  
        include_top=False,  
        weights=None,  
        input_shape=(32,32,3),  
    )  
  
    vgg_model.summary()
```

Model: "vgg19"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 32, 32, 3)]	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv4 (Conv2D)	(None, 8, 8, 256)	590080
=====		
block5_conv1 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv2 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv3 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv4 (Conv2D)	(None, 2, 2, 512)	2359808
block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0
=====		
Total params: 20,024,384		
Trainable params: 20,024,384		
Non-trainable params: 0		
=====		

```
In [24]: model = tf.keras.Sequential()  
        model.add(vgg_model)  
        model.add(Flatten())  
        model.add(Dense(1024, activation = 'relu'))  
        model.add(BatchNormalization())  
        model.add(Dense(1024, activation = 'relu'))  
        model.add(BatchNormalization())  
        model.add(Dense(256, activation = 'relu'))  
        model.add(BatchNormalization())  
        model.add(Dropout(0.5))  
        model.add(Dense(10, activation = 'softmax'))  
  
        model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
vgg19 (Functional)	(None, 1, 1, 512)	20024384

Layer (type)	Output Shape	Param #
vgg19 (Functional)	(None, 1, 1, 512)	20024384
flatten_1 (Flatten)	(None, 512)	0
dense_4 (Dense)	(None, 1024)	525312
batch_normalization_3 (Batch Normalization)	(None, 1024)	4096
dense_5 (Dense)	(None, 1024)	1049600
batch_normalization_4 (Batch Normalization)	(None, 1024)	4096
dense_6 (Dense)	(None, 256)	262400
batch_normalization_5 (Batch Normalization)	(None, 256)	1024
dropout_1 (Dropout)	(None, 256)	0
dense_7 (Dense)	(None, 10)	2570

=====  
 Total params: 21,873,482  
 Trainable params: 21,868,874  
 Non-trainable params: 4,608  
 =====

```
In [25]: optimizer = tf.keras.optimizers.SGD(learning_rate = 0.001, momentum = 0.9)
model.compile(optimizer= optimizer,
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

Trainable params: 21,868,874  
 Non-trainable params: 4,608

```
In [25]: optimizer = tf.keras.optimizers.SGD(learning_rate = 0.001, momentum = 0.9)
model.compile(optimizer= optimizer,
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```


```
In [75]: history = model.fit(
    train_datagen.flow(x_train, y_train, batch_size = 128),
    validation_data = val_datagen.flow(x_val, y_val, batch_size = 128),
    epochs = 100,
    verbose = 1,
    callbacks = [learning_rate_reduction]
)
```

```
Epoch 1/100
313/313 [=====] - 33s 106ms/step - loss: 1.0459 - accuracy: 0.6251 - val_loss: 1.1096 - val_accuracy: 0.6193 - lr: 1.0000e-05
Epoch 2/100
313/313 [=====] - 32s 102ms/step - loss: 1.0524 - accuracy: 0.6243 - val_loss: 1.1004 - val_accuracy: 0.6269 - lr: 1.0000e-05
Epoch 3/100
313/313 [=====] - 32s 101ms/step - loss: 1.0464 - accuracy: 0.6241 - val_loss: 1.0950 - val_accuracy: 0.6269 - lr: 1.0000e-05
Epoch 4/100
313/313 [=====] - 33s 105ms/step - loss: 1.0410 - accuracy: 0.6277 - val_loss: 1.1344 - val_accuracy: 0.6244 - lr: 1.0000e-05
Epoch 5/100
313/313 [=====] - 32s 102ms/step - loss: 1.0508 - accuracy: 0.6237 - val_loss: 1.0986 - val_accuracy: 0.6256 - lr: 1.0000e-05
Epoch 6/100
313/313 [=====] - 33s 104ms/step - loss: 1.0493 - accuracy: 0.6246 - val_loss: 1.0882 - val_accuracy: 0.6293 - lr: 1.0000e-05
```

```
In [76]: acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.figure()
plt.plot(acc,color = 'purple',label = 'Training Acuracy')
plt.plot(val_acc,color = 'blue',label = 'Validation Accuracy')
plt.legend()
```


Out[76]: <matplotlib.legend.Legend at 0x7f75101e8160>



```
In [77]: loss = history.history['loss']
val_loss = history.history['val_loss']

plt.figure()
plt.plot(loss,color = 'green',label = 'Training Loss')
plt.plot(val_loss,color = 'red',label = 'Validation Loss')
plt.legend()
```

Out[77]: <matplotlib.legend.Legend at 0x7f75101e8d30>



```
In [78]: X_test = tf.keras.applications.vgg19.preprocess_input(X_test)
y_pred = np.argmax(model.predict(X_test), axis=-1)
y_pred[:10]
```

```
313/313 [=====] - 3s 9ms/step
```

```
Out[78]: array([5, 1, 5, 5, 5, 5, 7, 5, 5, 7])
```

```
In [79]: from sklearn.metrics import confusion_matrix, accuracy_score
print('Testing Accuracy : ', accuracy_score(Y_test, y_pred))
```

```
Testing Accuracy : 0.1326
```

```
In [80]: cm = confusion_matrix(Y_test, y_pred)
cm
```

```
Out[80]: array([[ 7, 36,  0,  4,  0, 433,  0, 484,  0, 36],
 [61, 141,  0,  2,  1, 250,  0, 399,  0, 146],
 [ 1,  0,  0, 10,  0, 737,  0, 250,  0,  2],
 [ 0,  1,  0,  8,  3, 685,  0, 295,  0,  8],
 [ 0,  0,  0, 16,  3, 779,  0, 197,  0,  5],
 [ 1,  0,  0, 19,  2, 684,  0, 290,  0,  4],
 [ 3,  5,  0,  9,  1, 716,  0, 252,  0, 14],
 [ 1,  5,  0, 18,  2, 597,  0, 334,  0, 43],
 [15, 34,  0,  3,  0, 469,  0, 423,  0, 56],
 [27, 149,  0,  2,  1, 319,  0, 353,  0, 149]])
```

```
In [81]: import itertools
def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Greens):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
```

```
def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Greens):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=30)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    #print(cm)

    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```

```
In [82]: plt.figure(figsize=(8,8))
plot_confusion_matrix(cm, classes)
```

